

VonAI

FREE FRAMEWORK

The 30-Day Outcome Pilot

A 4-week sprint that turns "let's pilot AI" into a scoped workflow, a measured euro outcome, and a decision your CFO will sign.

von-ai.com · hello@von-ai.com

Why this exists

In 2023, "we're piloting AI" was a board-level answer. Leaders looked progressive. Vendors got paid. Nobody asked hard questions.

In 2026, that answer gets you fired.

Boards are asking 3 questions:

1. What did the pilot ship?
2. What did it save us in euros?
3. What's the plan to scale or kill it?

If you can't answer those, you don't have a pilot. You have a budget leak.

This framework gives you all 3 answers in 30 days.

The rules of the game

Before you start, agree on these. Skip them and the pilot drifts.

- 1. One workflow.** Not 5. Not "AI strategy." One concrete process where you can name the people, the tools, and the number.
- 2. One owner.** A single person whose name is on the outcome. Not a steering committee.
- 3. One number.** A baseline cost in euros and a target reduction. If you can't measure it, you can't pilot it.
- 4. 30 calendar days, not 30 working days.** The clock runs whether you're ready or not.
- 5. Smallest version wins.** Use tools you already pay for. No new platforms, no procurement cycles, no "we'll need to build it properly later."

The 4-week sprint

WEEK 1

Pick the workflow

Objective. Choose the one workflow worth piloting and write it down.

Steps

1. List the top 10 repetitive workflows in the team or business unit running the pilot.
2. Score each on 5 dimensions (1-5 scale): Frequency, Time cost, Error cost, AI fit, Build risk.
3. Multiply Frequency × Time × AI fit. Subtract Build risk. Top score wins.
4. Write a 1-page workflow brief: who runs it today, what triggers it, what output it produces, what it costs in euros per month.

OUTPUT

A 1-page workflow brief, signed by the owner.

TRAP TO AVOID

Picking the workflow that's most exciting instead of the one with the highest euros-per-month potential. Boring workflows ship. Exciting ones become "Phase 2."

Define the target

Objective. Lock the baseline cost and the target outcome before you build anything.

Steps

1. Calculate the baseline. Pick the metric: hours per week, errors per month, response time, throughput. Convert to euros where possible.
2. Set the target. A defensible v0 target is a 30% reduction or speed-up. Bigger numbers attract scrutiny. Smaller ones don't justify the work.
3. Define the measurement. Who measures, what counts, how often, what's the unit. Write it on the same page as the baseline.
4. Get the owner and the budget holder to sign the target sheet. If they won't sign, you don't have a pilot — you have a hobby project.

OUTPUT

A 1-page target sheet (baseline, target, measurement protocol, signatures).

TRAP TO AVOID

Soft targets like "improve efficiency" or "save time." If you can't put a euro number on it, the pilot won't survive the next budget review.

Build the smallest version

Objective. Get a working prototype in front of one real user, processing real cases, by the end of the week.

Steps

1. Use only tools you already have access to. An LLM, a spreadsheet, an automation tool, and whatever your business already pays for.
2. Build the prototype against 5 real, recent cases from the workflow. Not synthetic ones. Real ones with real outputs you can compare.
3. Test with 1 user. The user runs the prototype on their next 10 cases and logs what worked and what didn't.
4. Document the prompt, the agent, the automation, and the dependencies. Anyone with access should be able to re-run it.

OUTPUT

A working prototype + a usage log + a build doc.

TRAP TO AVOID

Spending Week 3 on tool selection or platform debates. The prototype is throwaway. The point is to learn what the workflow looks like when AI is in it, not to build something that scales on day 30.

Measure and decide

Objective. Run the prototype on real workflow for a full week, compare to baseline, and make a call.

Steps

1. Hand the prototype to the workflow owner. They run it for 5 working days against real cases.
2. Measure against the target sheet. Did you hit the 30% number? Came close? Missed entirely?
3. Compare 3 outcomes — Scale (hit or beat target, build the case for production), Iterate (within 50% of target, one more 2-week cycle on the bottleneck), Kill (no measurable signal, archive and move on).
4. Write the 1-page decision memo with "Scale / Iterate / Kill" in bold at the top. Send it to the budget holder before day 30.

OUTPUT

A 1-page decision memo.

TRAP TO AVOID

Treating "Iterate" as the default answer. If you're not sure, default to "Kill." Cheap to restart, expensive to drag out.

What you have at day 30

- A scoped workflow tied to a euro number
- A measured outcome
- A go / iterate / kill decision your CFO will sign
- A documented build that you (or a partner) can scale

That's the format boards fund. That's what replaces "we're piloting AI."

Templates included

The full framework comes with:

- Workflow scoring sheet (Excel / Google Sheets)
- Target sheet template (1 page)
- Build doc template (1 page)
- Decision memo template (1 page)

- Sample 30-day calendar with check-ins

Reply to the framework email and we'll send the spreadsheet pack.

Tempted to scope it bigger? Don't.

Bigger pilots fail. They fail because the scope grows faster than the team can ship, and because nobody can name the euro number when the budget holder asks.

Small pilots ship. Small pilots produce numbers. Numbers produce budget. Budget produces real builds.

The 30-day constraint isn't arbitrary. It's the longest you can run before the organization starts asking "what's happening with that AI thing?" Beat that question with a decision memo, not an apology.

When VonAI runs this for you

If you want help running the pilot, VonAI delivers it as a fixed-scope engagement. We pick the workflow with you, build the smallest version, measure it, and hand you the decision memo on day 30. Take the build in-house, or have us scale it.

Email: hello@von-ai.com · von-ai.com/ai-roi-sprint

VonAI helps small and medium service businesses turn AI from a budget item into a measurable line of revenue or savings.
von-ai.com