

The AI Operating Layer Checklist

The four things every business needs before scaling AI

You have agents. You have tools. What most businesses lack is the **operating layer** that makes them safe and reliable. Work through the four sections below and count your yeses.

01 Governance

- ☐ A written list of approved AI tools and approved use cases exists
- ☐ Every customer-facing AI output has a defined human checkpoint
- ☐ Employees have an AI usage policy they have actually read
- ☐ One named person is accountable for AI decisions

02 Memory

- ☐ You know what data each AI tool stores, and for how long
- ☐ There is a rule for what the system must forget, such as personal data or expired records
- ☐ Context carries across sessions where it should, and resets where it should not

03 Evaluation

- ☐ Each use case has a written "good enough to release" rule
- ☐ Someone checks outputs against that rule before a customer sees them
- ☐ You track failure rate, not just usage

04 Routing

- ☐ It is clear which task goes to which model or tool, and why
- ☐ High-risk tasks route to a stronger model or to a human
- ☐ There is a fallback when the primary tool fails or is uncertain

SCORE YOURSELF

0-5**Fragile**

You have random tools, not a system.

6-10**Partial**

One bad output away from a problem.

11-13**Operating**

The layer holds. Now you can scale.

START HERE

Do not fix everything at once. Pick your **highest-risk use case** and write its evaluation rule this week, in one sentence: "This output is good enough to release when _____".

Scored under 10? That gap is exactly what an AI ROI Sprint closes in the first session: governance, memory, evaluation and routing, written down and owned.

von-ai.com